

The World is Not Just Integers

❖ Programming languages support numbers with fraction

- ✧ Called **floating-point** numbers

- ✧ Examples:

3.14159265... (π)

2.71828... (e)

0.000000001 or 1.0×10^{-9} (seconds in a nanosecond)

86,400,000,000,000 or 8.64×10^{13} (nanoseconds in a day)

last number is a large integer that cannot fit in a 32-bit integer

❖ We use a **scientific notation** to represent

- ✧ Very small numbers (e.g. 1.0×10^{-9})

- ✧ Very large numbers (e.g. 8.64×10^{13})

- ✧ **Scientific notation**: $\pm d.f_1f_2f_3f_4 \dots \times 10^{\pm e_1e_2e_3}$

Floating-Point Numbers

❖ Examples of floating-point numbers in base 10 ...

✧ 5.341×10^3 , 0.05341×10^5 , -2.013×10^{-1} , -201.3×10^{-3}
↑ *decimal point*

❖ Examples of floating-point numbers in base 2 ...

✧ 1.00101×2^{23} , 0.0100101×2^{25} , -1.101101×2^{-3} , -1101.101×2^{-6}

✧ Exponents are kept in decimal for clarity

↑ *binary point*

✧ The binary number $(1101.101)_2 = 2^3 + 2^2 + 2^0 + 2^{-1} + 2^{-3} = 13.625$

❖ Floating-point numbers should be **normalized**

✧ Exactly **one non-zero digit** should appear **before the point**


- In a decimal number, this digit can be from **1 to 9**
- In a binary number, this digit should be **1**

✧ **Normalized FP Numbers:** 5.341×10^3 and -1.101101×2^{-3}

✧ **NOT Normalized:** 0.05341×10^5 and -1101.101×2^{-6}

Floating-Point Representation

- ❖ A floating-point number is represented by the triple
 - ✧ S is the **Sign bit** (0 is positive and 1 is negative)
 - Representation is called **sign and magnitude**
 - ✧ E is the **Exponent field** (signed)
 - Very large numbers have large positive exponents
 - Very small close-to-zero numbers have negative exponents
 - More bits in exponent field increases **range of values**
 - ✧ F is the **Fraction field** (fraction after binary point)
 - More bits in fraction field improves the **precision** of FP numbers



$$\text{Value of a floating-point number} = (-1)^S \times \text{val}(F) \times 2^{\text{val}(E)}$$

Next . . .

- ❖ Floating-Point Numbers
- ❖ IEEE 754 Floating-Point Standard
- ❖ Floating-Point Addition and Subtraction
- ❖ Floating-Point Multiplication
- ❖ MIPS Floating-Point Instructions

IEEE 754 Floating-Point Standard

- ❖ Found in virtually every computer invented since 1980
 - ✧ Simplified porting of floating-point numbers
 - ✧ Unified the development of floating-point algorithms
 - ✧ Increased the accuracy of floating-point numbers


❖ Single Precision Floating Point Numbers (32 bits)

- ✧ 1-bit sign + 8-bit exponent + 23-bit fraction



❖ Double Precision Floating Point Numbers (64 bits)

- ✧ 1-bit sign + 11-bit exponent + 52-bit fraction



Normalized Floating Point Numbers

❖ For a normalized floating point number (S , E , F)



❖ **Significand** is equal to $(1.F)_2 = (1.f_1 f_2 f_3 f_4 \dots)_2$

✧ IEEE 754 assumes hidden **1.** (**not stored**) for normalized numbers

✧ Significand is **1 bit longer** than fraction

❖ Value of a Normalized Floating Point Number is

$$(-1)^S \times (1.F)_2 \times 2^{\text{val}(E)}$$

$$(-1)^S \times (1.f_1 f_2 f_3 f_4 \dots)_2 \times 2^{\text{val}(E)}$$

$$(-1)^S \times (1 + f_1 \times 2^{-1} + f_2 \times 2^{-2} + f_3 \times 2^{-3} + f_4 \times 2^{-4} \dots)_2 \times 2^{\text{val}(E)}$$

$(-1)^S$ is 1 when S is 0 (positive), and -1 when S is 1 (negative)

Biased Exponent Representation

- ❖ How to represent a signed exponent? Choices are ...
 - ✧ Sign + magnitude representation for the exponent
 - ✧ Two's complement representation
 - ✧ Biased representation
- ❖ IEEE 754 uses **biased representation** for the **exponent**
 - ✧ Value of exponent = $\text{val}(E) = E - \text{Bias}$ (**Bias** is a constant)
- ❖ Recall that exponent field is **8 bits** for **single precision**
 - ✧ E can be in the range **0 to 255**
 - ✧ $E = 0$ and $E = 255$ are **reserved for special use** (discussed later)
 - ✧ $E = 1$ to **254** are used for **normalized** floating point numbers
 - ✧ **Bias = 127** (half of **254**), $\text{val}(E) = E - 127$
 - ✧ $\text{val}(E=1) = -126$, $\text{val}(E=127) = 0$, $\text{val}(E=254) = 127$

Biased Exponent - Cont'd

- ❖ For **double precision**, exponent field is **11 bits**
 - ✧ E can be in the range 0 to 2047
 - ✧ $E = 0$ and $E = 2047$ are **reserved for special use**
 - ✧ $E = 1$ to 2046 are used for **normalized** floating point numbers
 - ✧ Bias = 1023 (half of 2046), $\text{val}(E) = E - 1023$
 - ✧ $\text{val}(E=1) = -1022$, $\text{val}(E=1023) = 0$, $\text{val}(E=2046) = 1023$
- ❖ Value of a Normalized Floating Point Number is

$$(-1)^S \times (1.F)_2 \times 2^{E - \text{Bias}}$$

$$(-1)^S \times (1.f_1 f_2 f_3 f_4 \dots)_2 \times 2^{E - \text{Bias}}$$

$$(-1)^S \times (1 + f_1 \times 2^{-1} + f_2 \times 2^{-2} + f_3 \times 2^{-3} + f_4 \times 2^{-4} \dots)_2 \times 2^{E - \text{Bias}}$$

Examples of Single Precision Float

❖ What is the decimal value of this **Single Precision** float?

1	0	1	1	1	1	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

❖ **Solution:**

✧ Sign = 1 is negative

✧ Exponent = $(01111100)_2 = 124$, $E - \text{bias} = 124 - 127 = -3$

✧ Significand = $(1.0100 \dots 0)_2 = 1 + 2^{-2} = 1.25$ (**1.** is implicit)

✧ Value in decimal = $-1.25 \times 2^{-3} = -0.15625$

❖ What is the decimal value of?

0	1	0	0	0	0	1	0	0	1	0	0	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

❖ **Solution:**

implicit ↘

✧ Value in decimal = $+(1.01001100 \dots 0)_2 \times 2^{130-127} =$

$(1.01001100 \dots 0)_2 \times 2^3 = (1010.01100 \dots 0)_2 = 10.375$

Examples of Double Precision Float

❖ What is the decimal value of this **Double Precision** float ?

[illegible]

 **Solution:**

✧ Value of exponent = $(100000000101)_2 - \text{Bias} = 1029 - 1023 = 6$

✧ Value of double float = $(1.00101010 \dots 0)_2 \times 2^6$ (1. is implicit) = $(1001010.10 \dots 0)_2 = 74.5$

❖ What is the decimal value of ?

[illegible]

❖ **Do it yourself!** (answer should be $-1.5 \times 2^{-7} = -0.01171875$)

Converting FP Decimal to Binary

❖ Convert -0.8125 to binary in single and double precision

 **Solution:**

- ✧ Fraction bits can be obtained using multiplication by 2

- $0.8125 \times 2 = 1.625$
- $0.625 \times 2 = 1.25$
- $0.25 \times 2 = 0.5$
- $0.5 \times 2 = 1.0$
- Stop when fractional part is 0

$$0.8125 = (0.\textcolor{red}{1101})_2 = \frac{1}{2} + \frac{1}{4} + \frac{1}{16} = \frac{13}{16}$$

✧ Fraction = $(0.1101)_2 = (1.101)_2 \times 2^{-1}$ (Normalized)

✧ Exponent = $-1 + \text{Bias} = 126$ (single precision) and 1022 (double)

[illegible]

Single Precision

Double Precision

Largest Normalized Float

❖ What is the **Largest normalized float**?

❖ Solution for Single Precision:

[illegible]

✧ Exponent – bias = $254 - 127 = 127$ (largest exponent for SP)

✧ **Significand** = $(1.111 \dots 1)_2 = \text{almost } 2$

✧ Value in decimal $\approx 2 \times 2^{127} \approx 2^{128} \approx 3.4028 \dots \times 10^{38}$

❖ Solution for Double Precision:

[illegible]

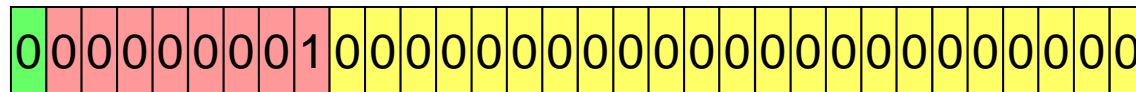
✧ Value in decimal $\approx 2 \times 2^{1023} \approx 2^{1024} \approx 1.79769 \dots \times 10^{308}$

❖ **Overflow:** exponent is **too large** to fit in the exponent field

Smallest Normalized Float

❖ What is the **smallest (in absolute value) normalized float**?

❖ **Solution for Single Precision:**

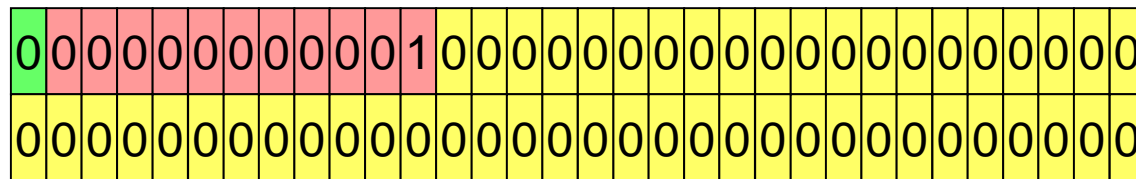


✧ Exponent – bias = $1 - 127 = -126$ (**smallest exponent for SP**)

✧ Significand = $(1.000 \dots 0)_2 = 1$

✧ Value in decimal = $1 \times 2^{-126} = 1.17549 \dots \times 10^{-38}$

❖ **Solution for Double Precision:**



✧ Value in decimal = $1 \times 2^{-1022} = 2.22507 \dots \times 10^{-308}$

❖ **Underflow:** exponent is **too small** to fit in exponent field

Zero, Infinity, and NaN

❖ Zero

- ✧ Exponent field $E = 0$ and fraction $F = 0$
- ✧ $+0$ and -0 are possible according to sign bit S

❖ Infinity

- ✧ Infinity is a special value represented with maximum E and $F = 0$
 - For **single precision** with 8-bit exponent: maximum $E = 255$
 - For **double precision** with 11-bit exponent: maximum $E = 2047$
- ✧ Infinity can result from overflow or division by zero
- ✧ $+\infty$ and $-\infty$ are possible according to sign bit S

❖ NaN (Not a Number)


- ✧ NaN is a special value represented with maximum E and $F \neq 0$
- ✧ Result from exceptional situations, such as $0/0$ or $\text{sqrt}(\text{negative})$
- ✧ Operation on a NaN results is NaN: $\text{Op}(X, \text{NaN}) = \text{NaN}$

Denormalized Numbers

- ❖ IEEE standard uses denormalized numbers to ...
 - ✧ Fill the gap between 0 and the smallest normalized float
 - ✧ Provide **gradual underflow** to zero
- ❖ **Denormalized**: exponent field E is 0 and fraction $F \neq 0$
 - ✧ Implicit **1.** before the fraction now becomes **0.** (**not normalized**)
- ❖ Value of denormalized number ($S, 0, F$)

Single precision: $(-1)^S \times (0.F)_2 \times 2^{-126}$

Double precision: $(-1)^S \times (0.F)_2 \times 2^{-1022}$



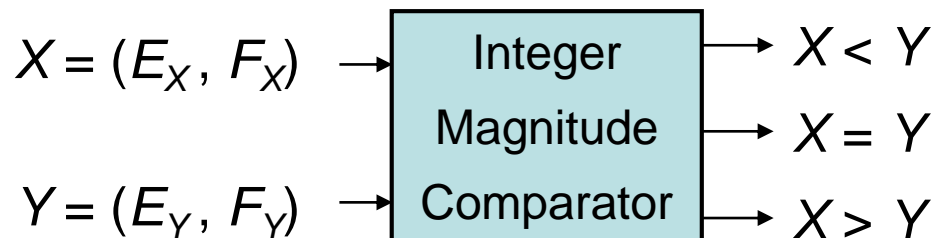
Summary of IEEE 754 Encoding

Single-Precision	Exponent = 8	Fraction = 23	Value
Normalized Number	1 to 254	Anything	$\pm (1.F)_2 \times 2^{E-127}$
Denormalized Number	0	nonzero	$\pm (0.F)_2 \times 2^{-126}$
Zero	0	0	± 0
Infinity	255	0	$\pm \infty$
NaN	255	nonzero	NaN

Double-Precision	Exponent = 11	Fraction = 52	Value
Normalized Number	1 to 2046	Anything	$\pm (1.F)_2 \times 2^{E-1023}$
Denormalized Number	0	nonzero	$\pm (0.F)_2 \times 2^{-1022}$
Zero	0	0	± 0
Infinity	2047	0	$\pm \infty$
NaN	2047	nonzero	NaN

Floating-Point Comparison

- ❖ IEEE 754 floating point numbers are ordered
 - ✧ Because exponent uses a biased representation ...
 - Exponent value and its binary representation have **same ordering**
 - ✧ Placing exponent before the fraction field **orders the magnitude**
 - **Larger exponent \Rightarrow larger magnitude**
 - **For equal exponents, Larger fraction \Rightarrow larger magnitude**
 - $0 < (0.F)_2 \times 2^{E_{\min}} < (1.F)_2 \times 2^{E-Bias} < \infty$ ($E_{\min} = 1 - Bias$)
 - ✧ Because sign bit is most significant \Rightarrow quick test of **signed <**
- ❖ Integer comparator can compare magnitudes



Next . . .

- ❖ Floating-Point Numbers
- ❖ IEEE 754 Floating-Point Standard
- ❖ Floating-Point Addition and Subtraction
- ❖ Floating-Point Multiplication
- ❖ MIPS Floating-Point Instructions

Floating Point Addition Example

❖ Consider Adding (Single-Precision Floating-Point):

$$+ 1.111001000000000000000010_2 \times 2^4$$

$$+ 1.1000000000000000110000101_2 \times 2^2$$

❖ Cannot add significands ... Why?

✧ Because **exponents are not equal**

❖ How to make exponents equal?

✧ **Shift the significand of the lesser exponent right**

✧ Difference between the two exponents = $4 - 2 = 2$

✧ So, **shift right** second number by **2** bits and increment exponent

$$1.1000000000000000110000101_2 \times 2^2$$

$$= 0.011000000000000001100001 \ 01_2 \times 2^4$$

Floating-Point Addition - cont'd

❖ Now, **ADD the Significands:**

$$\begin{array}{rcl}
 + 1.111001000000000000000010 & \times 2^4 & \\
 + 1.1000000000000000110000101 & \times 2^2 & \\
 \hline
 + 1.111001000000000000000010 & \times 2^4 & \\
 + 0.0110000000000000001100001\ 01 & \times 2^4 \text{ (shift right)} & \\
 \hline
 + \mathbf{10.01000100000000001100011\ 01} & \times 2^4 \text{ (result)} &
 \end{array}$$

❖ Addition produces a **carry bit**, result is NOT normalized

❖ **Normalize Result (shift right and increment exponent):**

$$\begin{array}{rcl}
 + \mathbf{10.01000100000000001100011\ 01} & \times 2^4 & \\
 = + \mathbf{1.00100010000000000110001\ 101} & \times 2^5 &
 \end{array}$$

Rounding

- ❖ Single-precision requires only 23 fraction bits
- ❖ However, Normalized result can contain additional bits

$$1.0010001000000000000110001 \mid \overset{\text{Round Bit: } R=1}{\textcircled{1}} \overset{\text{Sticky Bit: } S=1}{\textcircled{01}} \times 2^5$$

- ❖ Two extra bits are needed for rounding
 - ✧ Round bit: appears just after the normalized result
 - ✧ Sticky bit: appears after the round bit (OR of all additional bits)
- ❖ Since **RS = 11**, increment fraction to round to nearest

$$1.0010001000000000000110001 \times 2^5$$

$$+1$$

$$1.00100010000000000001100\mathbf{10} \times 2^5 \text{ (Rounded)}$$

Floating-Point Subtraction Example

❖ Sometimes, addition is converted into subtraction

✧ If the sign bits of the operands are different

❖ Consider Adding:

+ 1.00000000101100010001101 × 2⁻⁶

- 1.000000000000000010011010 × 2⁻¹

+ 0.00001000000001011000100 01101 × 2⁻¹ (shift right 5 bits)

- 1.000000000000000010011010 × 2⁻¹

0 0.00001000000001011000100 01101 × 2⁻¹

1 0.111111111111111101100110 × 2⁻¹ (2's complement)

1 1.00001000000001000101010 01101 × 2⁻¹ (ADD)

- 0.111101111111110111010101 10011 × 2⁻¹ (2's complement)

❖ 2's complement of result is required if result is negative

Floating-Point Subtraction - cont'd

$$\begin{array}{r}
 + 1.000000000101100010001101 \times 2^{-6} \\
 - 1.0000000000000000010011010 \times 2^{-1} \\
 \hline
 - 0.111101111111110111010101 \ 10011 \times 2^{-1} \text{ (result is negative)}
 \end{array}$$

❖ Result should be normalized

- ✧ For subtraction, we can have **leading zeros**. To normalize, count the number of leading zeros, then shift result left and decrement the exponent accordingly.

$$\begin{array}{r}
 - 0.111101111111110111010101 \overset{\text{Guard bit}}{\text{(1)}} 0011 \times 2^{-1} \\
 \hline
 - 1.111011111111110111010101 \text{1} \downarrow 0011 \times 2^{-2} \text{ (Normalized)}
 \end{array}$$

❖ **Guard bit**: guards against loss of a fraction bit

- ✧ Needed for subtraction, when result has a leading zero and should be normalized.

Floating-Point Subtraction - cont'd

- ❖ Next, normalized result should be **rounded**

$$\begin{array}{r}
 \text{Guard bit} \\
 - \quad 0.111101111111110111010101 \quad \textcircled{1} \quad 0 \quad 011 \times 2^{-1} \\
 \hline
 - \quad 1.111011111111110111010101 \quad \textcircled{1} \quad \textcircled{0} \quad \textcircled{011} \times 2^{-2} \text{ (Normalized)} \\
 \text{Round bit: } R=0 \quad \text{Sticky bit: } S=1
 \end{array}$$

- ❖ Since **R = 0**, it is more accurate to **truncate** the result even if **S = 1**. We simply discard the extra bits.

$$\begin{array}{r}
 - \quad 1.1110111111111101110101011 \quad 0 \quad 011 \times 2^{-2} \text{ (Normalized)} \\
 \hline
 - \quad 1.1110111111111101110101011 \quad \times 2^{-2} \text{ (Rounded to nearest)}
 \end{array}$$

- ❖ IEEE 754 Representation of Result

1	0	1	1	1	1	0	1	1	1	1	0	1	1	1	1	1	1	0	1	1	1	0	1	0	1	0	1	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Rounding to Nearest Even

- ❖ Normalized result has the form: $1.f_1 f_2 \dots f_l \mathbf{R} \mathbf{S}$
 - ✧ The **round bit R** appears after the last fraction bit f_l
 - ✧ The **sticky bit S** is the OR of all remaining additional bits
- ❖ **Round to Nearest Even**: default rounding mode
- ❖ Four cases for **RS**:
 - ✧ **RS = 00** → Result is Exact, no need for rounding
 - ✧ **RS = 01** → **Truncate** result by discarding **RS**
 - ✧ **RS = 11** → **Increment** result: ADD 1 to last fraction bit
 - ✧ **RS = 10** → Tie Case (either truncate or increment result)
 - Check Last fraction bit f_l (f_{23} for single-precision or f_{52} for double)
 - If f_l is **0** then **truncate** result to keep fraction even
 - If f_l is **1** then **increment** result to make fraction even

Additional Rounding Modes

❖ IEEE 754 standard specifies four rounding modes:

1. **Round to Nearest Even**: described in previous slide

2. **Round toward +Infinity**: result is rounded up

Increment result if **sign is positive and R or S = 1**

3. **Round toward -Infinity**: result is rounded down

Increment result if **sign is negative and R or S = 1**

4. **Round toward 0**: always truncate result

❖ Rounding or Incrementing result might generate a carry

✧ This occurs when all fraction bits are **1**

✧ Re-Normalize after Rounding step is required only in this case

Example on Rounding

❖ Round following result using IEEE 754 rounding modes:

$$-1.11111111111111111111111111111111 \overset{\text{Round Bit}}{\underset{\uparrow}{(1)}} \overset{\text{Sticky Bit}}{\underset{\uparrow}{(0)}} \times 2^{-7}$$

❖ Round to Nearest Even: *Round Bit* \uparrow \uparrow *Sticky Bit*

✧ Increment result since **RS = 10 and $f_{23} = 1$**

✧ Incremented result: **-10.000000000000000000000000000000** $\times 2^{-7}$

✧ Renormalize and increment exponent (**because of carry**)

✧ Final rounded result: **-1.000000000000000000000000000000** $\times 2^{-6}$

❖ Round towards $+\infty$: **Truncate** result since **negative**

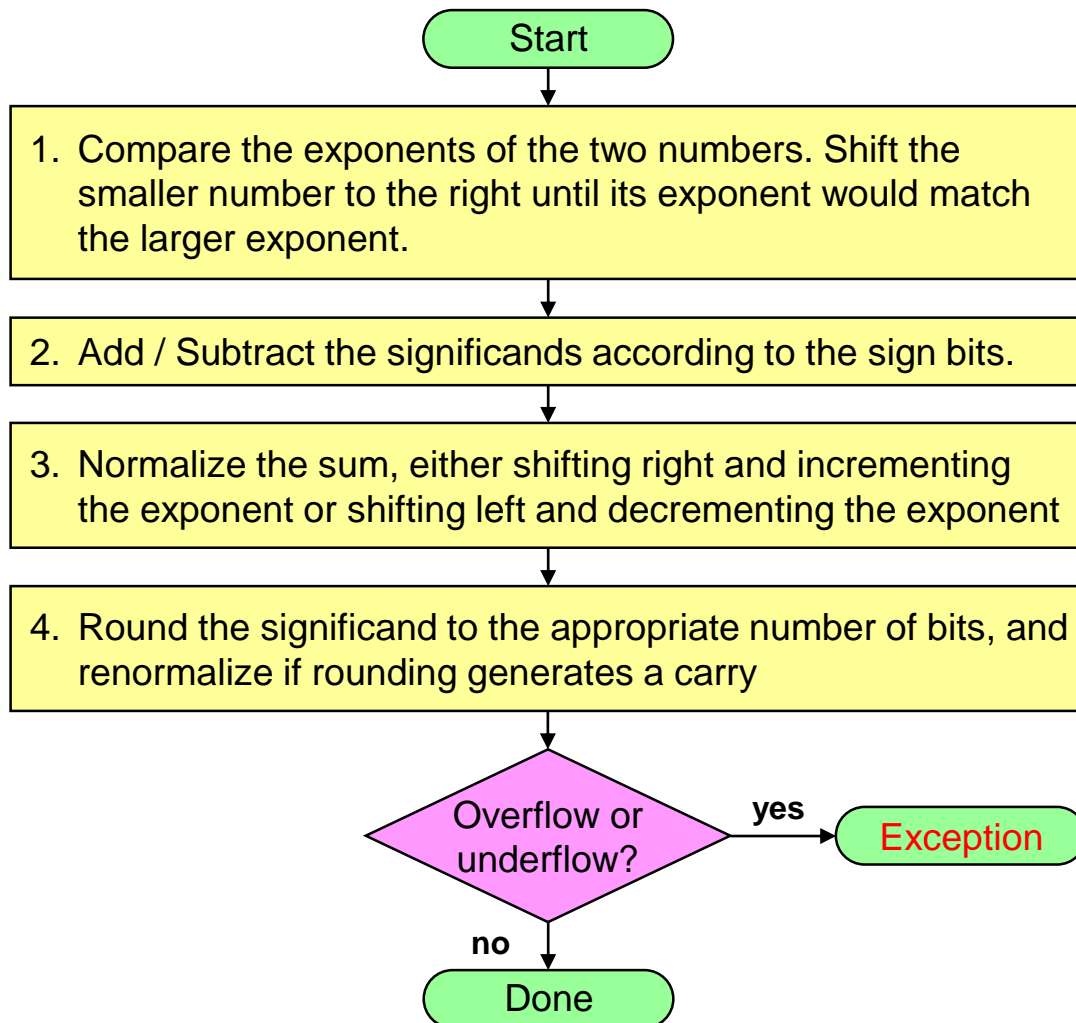
✧ Truncated Result: **-1.111111111111111111111111111111** $\times 2^{-7}$

❖ Round towards $-\infty$: **Increment** since **negative and $R = 1$**

✧ Final rounded result: **-1.000000000000000000000000000000** $\times 2^{-6}$

❖ Round towards 0: **Truncate always**

Floating Point Addition / Subtraction



Shift significand right by

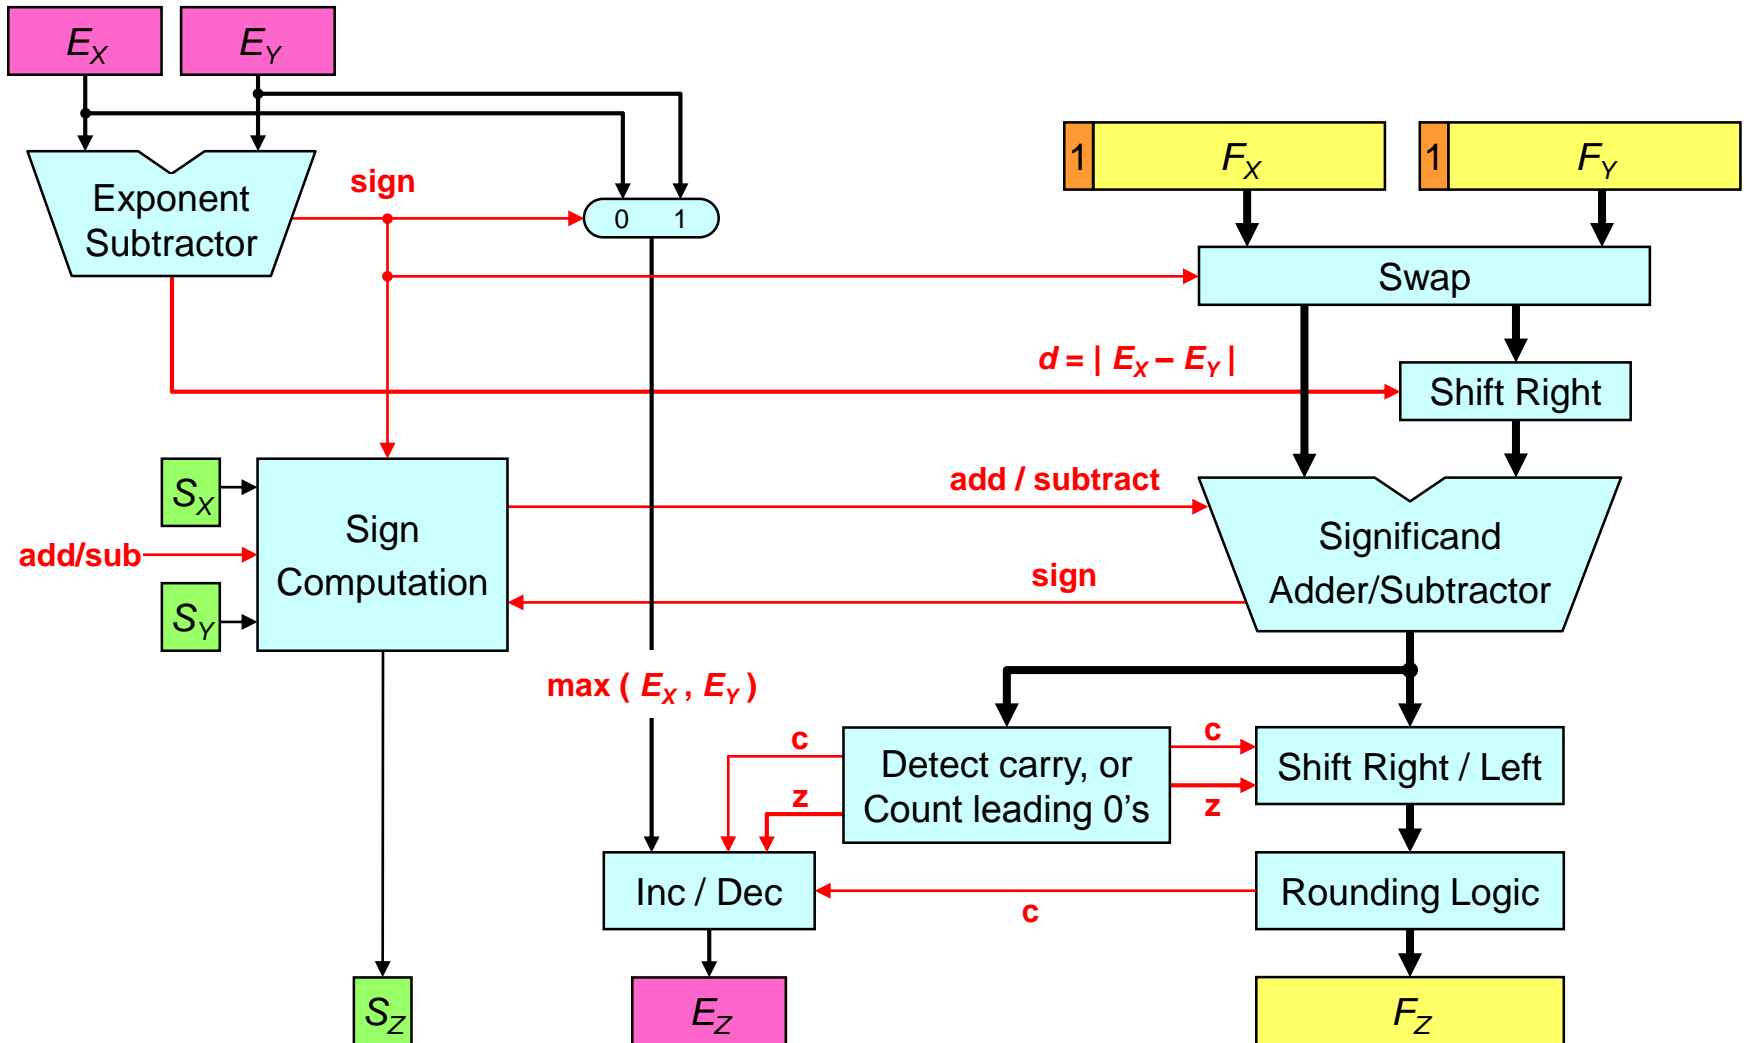
$$d = |E_X - E_Y|$$

Add significands when signs of X and Y are identical,
Subtract when different
 $X - Y$ becomes $X + (-Y)$

Normalization shifts right by 1 if there is a carry, or shifts left by the number of leading zeros in the case of subtraction

Rounding either truncates fraction, or adds a 1 to least significant fraction bit

Floating Point Adder Block Diagram



Next . . .

- ❖ Floating-Point Numbers
- ❖ IEEE 754 Floating-Point Standard
- ❖ Floating-Point Addition and Subtraction
- ❖ Floating-Point Multiplication
- ❖ MIPS Floating-Point Instructions

Floating Point Multiplication Example

❖ Consider multiplying:

$$\begin{array}{r} -1.110 \ 1000 \ 0100 \ 0000 \ 1010 \ 0001_2 \times 2^{-4} \\ \times \ 1.100 \ 0000 \ 0001 \ 0000 \ 0000 \ 0000_2 \times 2^{-2} \end{array}$$

❖ Unlike addition, we **add the exponents** of the operands

✧ Result exponent value = $(-4) + (-2) = -6$

❖ Using the biased representation: $E_Z = E_X + E_Y - \text{Bias}$

✧ $E_X = (-4) + 127 = 123$ (**Bias = 127** for single precision)

✧ $E_Y = (-2) + 127 = 125$

✧ $E_Z = 123 + 125 - 127 = 121$ (**value = -6**)

❖ Sign bit of product can be computed independently

❖ Sign bit of product = $\text{Sign}_X \mathbf{XOR} \text{Sign}_Y = \mathbf{1}$ (**negative**)

Floating-Point Multiplication, cont'd

❖ Now multiply the significands:

$$\begin{array}{r} \text{(Multiplicand)} \quad 1.11010000100000010100001 \\ \text{(Multiplier)} \quad \times \quad 1.10000000001000000000000 \\ \hline 111010000100000010100001 \\ 111010000100000010100001 \\ 1.11010000100000010100001 \\ \hline 10.1011100011111011111100110010100001000000000000 \end{array}$$

❖ 24 bits \times 24 bits \rightarrow 48 bits (double number of bits)

❖ Multiplicand $\times 0 = 0$ Zero rows are eliminated

❖ Multiplicand $\times 1 =$ Multiplicand (shifted left)

Floating-Point Multiplication, cont'd

❖ Normalize Product:

$$-10.10111000111110111111001100\dots \times 2^{-6}$$

Shift right and increment exponent because of **carry bit**

$$= -1.010111000111110111111001100\dots \times 2^{-5}$$

❖ Round to Nearest Even: (keep only 23 fraction bits)

$$1.01011100011111011111100 \mid \boxed{1} \boxed{100\dots} \times 2^{-5}$$

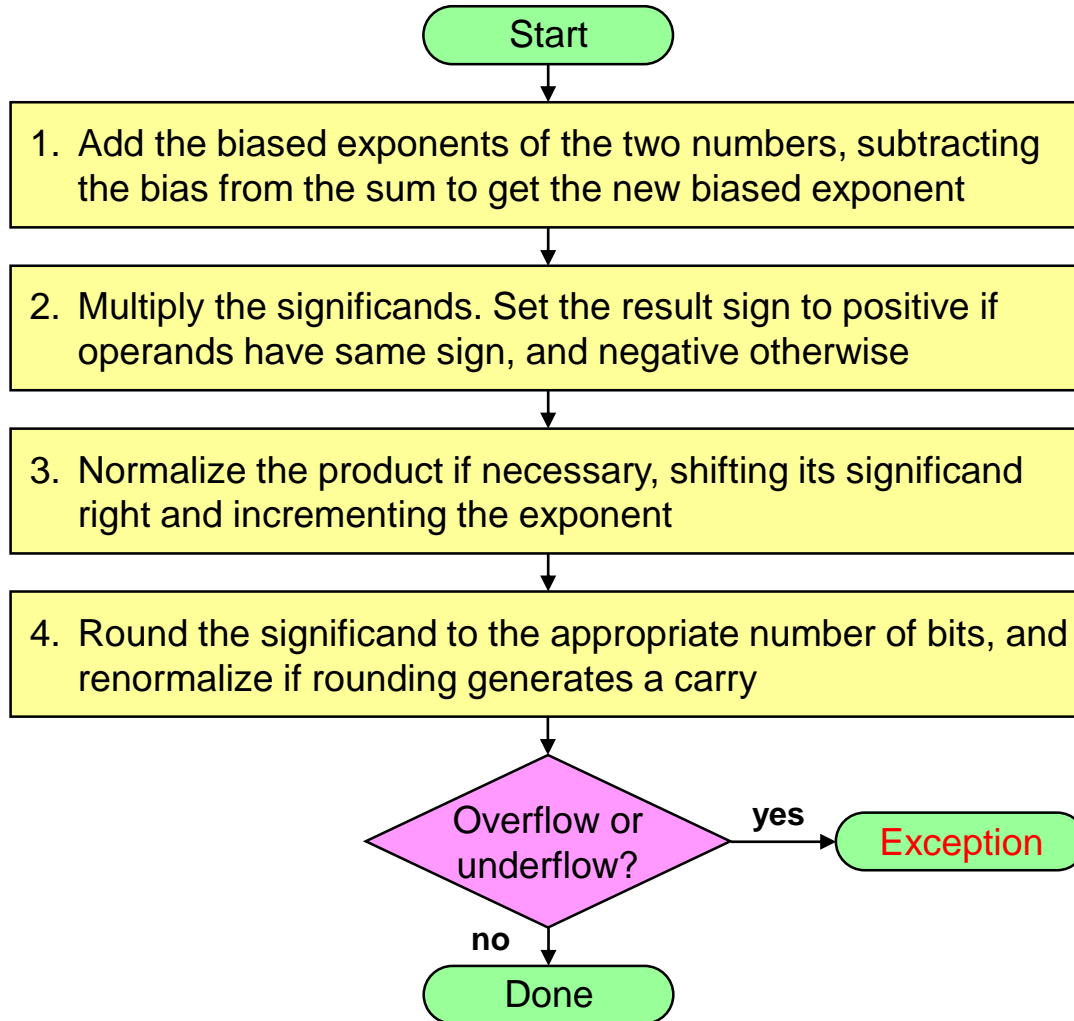
Round bit = **1**, Sticky bit = **1**, so increment fraction

$$\text{Final result} = -1.01011100011111011111101 \times 2^{-5}$$

❖ IEEE 754 Representation

1	0	1	1	1	1	0	1	0	0	1	0	1	1	1	0	0	0	1	1	1	1	1	0	1	1	1	1	1	0	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Floating Point Multiplication



Biased Exponent Addition

$$E_Z = E_X + E_Y - Bias$$

Result sign $S_Z = S_X \text{ xor } S_Y$ can be computed independently

Since the operand significands $1.F_X$ and $1.F_Y$ are ≥ 1 and < 2 , their product is ≥ 1 and < 4 .

To normalize product, we need to shift right at most by 1 bit and increment exponent

Rounding either truncates fraction, or adds a 1 to least significant fraction bit

Extra Bits to Maintain Precision

- ❖ Floating-point numbers are approximations for ...
 - ✧ Real numbers that they cannot represent
- ❖ Infinite variety of real numbers exist between 1.0 and 2.0
 - ✧ However, exactly 2^{23} fractions represented in Single Precision
 - ✧ Exactly 2^{52} fractions can be represented in Double Precision
- ❖ Extra bits are generated in intermediate results when ...
 - ✧ Shifting and adding/subtracting a p -bit significand
 - ✧ Multiplying two p -bit significands (product is $2p$ bits)
- ❖ But when packing result fraction, extra bits are discarded
- ❖ Few extra bits are needed: guard, round, and sticky bits
- ❖ Minimize hardware but without compromising accuracy

Advantages of IEEE 754 Standard

- ❖ Used predominantly by the industry
- ❖ Encoding of exponent and fraction simplifies comparison
 - ✧ Integer comparator used to compare magnitude of FP numbers
- ❖ Includes special exceptional values: NaN and $\pm\infty$
 - ✧ Special rules are used such as:
 - $0/0$ is NaN, $\text{sqrt}(-1)$ is NaN, $1/0$ is ∞ , and $1/\infty$ is 0
 - ✧ Computation may continue in the face of exceptional conditions
- ❖ Denormalized numbers to fill the gap
 - ✧ Between smallest normalized number $1.0 \times 2^{E_{min}}$ and zero
 - ✧ Denormalized numbers, values $0.F \times 2^{E_{min}}$, are closer to zero
 - ✧ Gradual underflow to zero

Floating Point Complexities

- ❖ Operations are somewhat more complicated
- ❖ In addition to **overflow** we can have **underflow**
- ❖ Accuracy can be a big problem
 - ✧ Extra bits to maintain precision: **guard**, **round**, and **sticky**
 - ✧ Four **rounding modes**
 - ✧ Division by zero yields **Infinity**
 - ✧ Zero divide by zero yields **Not-a-Number**
 - ✧ Other complexities
- ❖ Implementing the standard can be tricky
 - ✧ See text for description of 80x86 and Pentium bug!
- ❖ Not using the standard can be even worse

Accuracy can be a Big Problem

Value1	Value2	Value3	Value4	Sum
1.0E+30	-1.0E+30	9.5	-2.3	7.2
1.0E+30	9.5	-1.0E+30	-2.3	-2.3
1.0E+30	9.5	-2.3	-1.0E+30	0

- ❖ Adding double-precision floating-point numbers (Excel)
- ❖ Floating-Point addition is NOT associative
- ❖ Produces different sums for the same data values
- ❖ Rounding errors when the difference in exponent is large

Next . . .

- ❖ Floating-Point Numbers
- ❖ IEEE 754 Floating-Point Standard
- ❖ Floating-Point Addition and Subtraction
- ❖ Floating-Point Multiplication
- ❖ MIPS Floating-Point Instructions

MIPS Floating Point Coprocessor

- ❖ Called **Coprocessor 1** or the **Floating Point Unit (FPU)**
- ❖ 32 separate floating point registers: \$f0, \$f1, ..., \$f31
- ❖ FP registers are 32 bits for single precision numbers
- ❖ Even-odd register pair form a double precision register
- ❖ Use the even number for double precision registers
 - ✧ \$f0, \$f2, \$f4, ..., \$f30 are used for double precision
- ❖ Separate FP instructions for single/double precision
 - ✧ Single precision: **add.s, sub.s, mul.s, div.s** (**.s extension**)
 - ✧ Double precision: **add.d, sub.d, mul.d, div.d** (**.d extension**)
- ❖ FP instructions are more complex than the integer ones
 - ✧ Take more cycles to execute

FP Arithmetic Instructions

Instruction		Meaning	Format					
add.s	fd, fs, ft	$(fd) = (fs) + (ft)$	0x11	0	ft ⁵	fs ⁵	fd ⁵	0
add.d	fd, fs, ft	$(fd) = (fs) + (ft)$	0x11	1	ft ⁵	fs ⁵	fd ⁵	0
sub.s	fd, fs, ft	$(fd) = (fs) - (ft)$	0x11	0	ft ⁵	fs ⁵	fd ⁵	1
sub.d	fd, fs, ft	$(fd) = (fs) - (ft)$	0x11	1	ft ⁵	fs ⁵	fd ⁵	1
mul.s	fd, fs, ft	$(fd) = (fs) \times (ft)$	0x11	0	ft ⁵	fs ⁵	fd ⁵	2
mul.d	fd, fs, ft	$(fd) = (fs) \times (ft)$	0x11	1	ft ⁵	fs ⁵	fd ⁵	2
div.s	fd, fs, ft	$(fd) = (fs) / (ft)$	0x11	0	ft ⁵	fs ⁵	fd ⁵	3
div.d	fd, fs, ft	$(fd) = (fs) / (ft)$	0x11	1	ft ⁵	fs ⁵	fd ⁵	3
sqrt.s	fd, fs	$(fd) = \text{sqrt}(fs)$	0x11	0	0	fs ⁵	fd ⁵	4
sqrt.d	fd, fs	$(fd) = \text{sqrt}(fs)$	0x11	1	0	fs ⁵	fd ⁵	4
abs.s	fd, fs	$(fd) = \text{abs}(fs)$	0x11	0	0	fs ⁵	fd ⁵	5
abs.d	fd, fs	$(fd) = \text{abs}(fs)$	0x11	1	0	fs ⁵	fd ⁵	5
neg.s	fd, fs	$(fd) = -(fs)$	0x11	0	0	fs ⁵	fd ⁵	7
neg.d	fd, fs	$(fd) = -(fs)$	0x11	1	0	fs ⁵	fd ⁵	7

FP Load/Store Instructions

❖ Separate floating point load/store instructions

- ✧ `lwc1`: load word coprocessor 1
- ✧ `ldc1`: load double coprocessor 1
- ✧ `swc1`: store word coprocessor 1
- ✧ `sdcl`: store double coprocessor 1

General purpose register is used as the **base** register

Instruction		Meaning	Format			
<code>lwc1</code>	<code>\$f2, 40(\$t0)</code>	$(\$f2) = \text{Mem}[(\$t0)+40]$	0x31	<code>\$t0</code>	<code>\$f2</code>	$\text{im}^{16} = 40$
<code>ldc1</code>	<code>\$f2, 40(\$t0)</code>	$(\$f2) = \text{Mem}[(\$t0)+40]$	0x35	<code>\$t0</code>	<code>\$f2</code>	$\text{im}^{16} = 40$
<code>swc1</code>	<code>\$f2, 40(\$t0)</code>	$\text{Mem}[(\$t0)+40] = (\$f2)$	0x39	<code>\$t0</code>	<code>\$f2</code>	$\text{im}^{16} = 40$
<code>sdcl</code>	<code>\$f2, 40(\$t0)</code>	$\text{Mem}[(\$t0)+40] = (\$f2)$	0x3d	<code>\$t0</code>	<code>\$f2</code>	$\text{im}^{16} = 40$

❖ Better names can be used for the above instructions

- ✧ `l.s` = `lwc1` (load FP single), `l.d` = `ldc1` (load FP double)
- ✧ `s.s` = `swc1` (store FP single), `s.d` = `sdcl` (store FP double)

FP Data Movement Instructions

- ❖ Moving data between general purpose and FP registers
 - ✧ **mfc1**: move from coprocessor 1 (to general purpose register)
 - ✧ **mtc1**: move to coprocessor 1 (from general purpose register)
- ❖ Moving data between FP registers
 - ✧ **mov.s**: move single precision float
 - ✧ **mov.d**: move double precision float = even/odd pair of registers

Instruction		Meaning	Format					
mfc1	\$t0, \$f2	(\$t0) = (\$f2)	0x11	0	\$t0	\$f2	0	0
mtc1	\$t0, \$f2	(\$f2) = (\$t0)	0x11	4	\$t0	\$f2	0	0
mov.s	\$f4, \$f2	(\$f4) = (\$f2)	0x11	0	0	\$f2	\$f4	6
mov.d	\$f4, \$f2	(\$f4) = (\$f2)	0x11	1	0	\$f2	\$f4	6

FP Convert Instructions

❖ Convert instruction: `cvt.x.y`

✧ Convert to **destination** format **x** from **source** format **y**

❖ Supported formats

- ✧ Single precision float = **.s** (single precision float in FP register)
- ✧ Double precision float = **.d** (double float in even-odd FP register)
- ✧ Signed integer word = **.w** (signed integer in FP register)

Instruction	Meaning	Format						
<code>cvt.s.w fd, fs</code>	to single from integer	0x11	0	0	fs ⁵	fd ⁵	0x20	
<code>cvt.s.d fd, fs</code>	to single from double	0x11	1	0	fs ⁵	fd ⁵	0x20	
<code>cvt.d.w fd, fs</code>	to double from integer	0x11	0	0	fs ⁵	fd ⁵	0x21	
<code>cvt.d.s fd, fs</code>	to double from single	0x11	1	0	fs ⁵	fd ⁵	0x21	
<code>cvt.w.s fd, fs</code>	to integer from single	0x11	0	0	fs ⁵	fd ⁵	0x24	
<code>cvt.w.d fd, fs</code>	to integer from double	0x11	1	0	fs ⁵	fd ⁵	0x24	

FP Compare and Branch Instructions

- ❖ FP unit (co-processor 1) has a condition flag
 - ✧ Set to 0 (false) or 1 (true) by any comparison instruction
- ❖ Three comparisons: equal, less than, less than or equal
- ❖ Two branch instructions based on the condition flag

Instruction		Meaning	Format					
c.eq.s	fs, ft	cflag = ((fs) == (ft))	0x11	0	ft ⁵	fs ⁵	0	0x32
c.eq.d	fs, ft	cflag = ((fs) == (ft))	0x11	1	ft ⁵	fs ⁵	0	0x32
c.lt.s	fs, ft	cflag = ((fs) < (ft))	0x11	0	ft ⁵	fs ⁵	0	0x3c
c.lt.d	fs, ft	cflag = ((fs) < (ft))	0x11	1	ft ⁵	fs ⁵	0	0x3c
c.le.s	fs, ft	cflag = ((fs) <= (ft))	0x11	0	ft ⁵	fs ⁵	0	0x3e
c.le.d	fs, ft	cflag = ((fs) <= (ft))	0x11	1	ft ⁵	fs ⁵	0	0x3e
bc1f	Label	branch if (cflag == 0)	0x11	8	0	im ¹⁶		
bc1t	Label	branch if (cflag == 1)	0x11	8	1	im ¹⁶		

Example 1: Area of a Circle

```
.data
    pi:      .double      3.1415926535897924
    msg:     .asciiz      "Circle Area = "
.text
main:
    ldc1     $f2, pi      # $f2,3 = pi
    li       $v0, 7       # read double (radius)
    syscall  # $f0,1 = radius
    mul.d    $f12, $f0, $f0 # $f12,13 = radius*radius
    mul.d    $f12, $f2, $f12 # $f12,13 = area
    la       $a0, msg
    li       $v0, 4       # print string (msg)
    syscall
    li       $v0, 3       # print double (area)
    syscall # print $f12,13
```

Example 2: Matrix Multiplication

```
void mm (int n, double x[n][n], y[n][n], z[n][n]) {  
    for (int i=0; i!=n; i=i+1)  
        for (int j=0; j!=n; j=j+1) {  
            double sum = 0.0;  
            for (int k=0; k!=n; k=k+1)  
                sum = sum + y[i][k] * z[k][j];  
            x[i][j] = sum;  
        }  
}
```


- ❖ Matrices **x**, **y**, and **z** are **n×n double precision** float
- ❖ Matrix size is passed in **\$a0 = n**
- ❖ Array addresses are passed in **\$a1**, **\$a2**, and **\$a3**
- ❖ What is the MIPS assembly code for the procedure?

Address Calculation for 2D Arrays

❖ Row-Major Order: 2D arrays are stored as rows

❖ Calculate Address of: $X[i][j]$

= Address of $X + (i \times n + j) \times 8$ (8 bytes per element)



❖ Address of $Y[i][k] = \text{Address of } Y + (i \times n + k) \times 8$

❖ Address of $Z[k][j] = \text{Address of } Z + (k \times n + j) \times 8$

Matrix Multiplication Procedure - 1/3

❖ Initialize Loop Variables

```
mm:  addu    $t1, $0, $0      # $t1 = i = 0; for 1st loop
L1:  addu    $t2, $0, $0      # $t2 = j = 0; for 2nd loop
L2:  addu    $t3, $0, $0      # $t3 = k = 0; for 3rd loop
      sub.d  $f0, $f0, $f0    # $f0 = sum = 0.0
```

❖ Calculate address of $y[i][k]$ and load it into $\$f2, \$f3$

❖ Skip i rows ($i \times n$) and add k elements

```
L3:  mul     $t4, $t1, $a0    # $t4 = i*size(row) = i*n
      addu   $t4, $t4, $t3    # $t4 = i*n + k
      sll    $t4, $t4, 3      # $t4 = (i*n + k)*8
      addu   $t4, $a2, $t4    # $t4 = address of y[i][k]
      l.d    $f2, 0($t4)     # $f2 = y[i][k]
```

Matrix Multiplication Procedure - 2/3

- ❖ Similarly, calculate address and load value of $z[k][j]$
- ❖ Skip k rows ($k \times n$) and add j elements

```
mul    $t5, $t3, $a0    # $t5 = k*size(row) = k*n
addu   $t5, $t5, $t2     # $t5 = k*n + j
sll    $t5, $t5, 3       # $t5 = (k*n + j)*8
addu   $t5, $a3, $t5     # $t5 = address of z[k][j]
ld     $f4, 0($t5)       # $f4 = z[k][j]
```

- ❖ Now, multiply $y[i][k]$ by $z[k][j]$ and add it to $\$f0$

```
mul.d  $f6, $f2, $f4     # $f6 = y[i][k]*z[k][j]
add.d  $f0, $f0, $f6     # $f0 = sum
addiu  $t3, $t3, 1       # k = k + 1
bne    $t3, $a0, L3      # loop back if (k != n)
```

Matrix Multiplication Procedure - 3/3

❖ Calculate address of $x[i][j]$ and store sum

```
mul    $t6, $t1, $a0    # $t6 = i*size(row) = i*n
addu   $t6, $t6, $t2    # $t6 = i*n + j
sll    $t6, $t6, 3      # $t6 = (i*n + j)*8
addu   $t6, $a1, $t6    # $t6 = address of x[i][j]
s.d    $f0, 0($t6)     # x[i][j] = sum
```

❖ Repeat outer loops: L2 (for $j = \dots$) and L1 (for $i = \dots$)

```
addiu  $t2, $t2, 1      # j = j + 1
bne    $t2, $a0, L2     # loop L2 if (j != n)
addiu  $t1, $t1, 1      # i = i + 1
bne    $t1, $a0, L1     # loop L1 if (i != n)
```

❖ Return:

```
jr     $ra              # return
```